

Architectures for massive data management

Key-value stores and Redis

Ioana Manolescu

INRIA Saclay & Ecole Polytechnique

ioana.manolescu@inria.fr

<http://pages.saclay.inria.fr/ioana.manolescu/>

M2 Data and Knowledge
Université de Paris Saclay

Architectures for Massive DM
D&K / UPSay 2016-2017

Ioana Manolescu

1

Key-value stores


- Relatively recent class of systems, developed as part of the « NoSQL » movement
- Main idea:
 - Trade simplicity for speed and scale
- Extremely simple data model
 - **key**=short byte sequence / integer
 - **value**=byte sequence (may recognize integers)
- No QL. Operations: **PUT(k, v)** and **GET(k,v)**
- **ACID** properties depending on the system; at least atomic PUT and GET
 - Some are in-memory thus no durability at all

Key-value data models

- Simplest model:
 - One key – one value
- Extensions:
 - **Organization:** key-value pairs belong to « collections » or « databases » or « tables »
 - **Multiplicity:** set or list of values
 - **Internal tructure:**
 - One key – a list of *attributes*
 - Each attribute has a *name* and a *value / set of values*

Sample key-value data model: DynamoDB

- Provided by Amazon Web Services (AWS)



- Naming may vary (there is no standard). See doc.
- Although it is called « table », *items in the same table may have nothing in common!*
- The interface is very similar to the so-called « Big Tables » (to be seen)

Redis: one of the most popular key-value stores

- Data model:
 - Hash (a set of key-value pairs on the same key)
 - List
 - Set
 - Values cannot be lists nor sets (no nesting!)
 - Databases
- Operations
 - Put, get
 - Set operations (union, intersection)
 - List operations: left/right push/pop (→queue / stack)
 - Arithmetic operations (attempts type conversion to integers)

Redis: first lab

- Install Redis, launch it
- Follow the tutorial
- Write in Redis a toy application of a database of books which can be borrowed in a library.
 - All books have an ISBN, a title and an author.
 - Books may also have other properties, e.g. language, publication year, edition...
 - Make books expire after a while (if no one borrows the book)
 - Make a Redis client subscribe to books by a certain author
 - With the help of a programming language interface to Redis, index a newly published item by the words in its description. Make a Redis client subscribe to items having certain keywords.