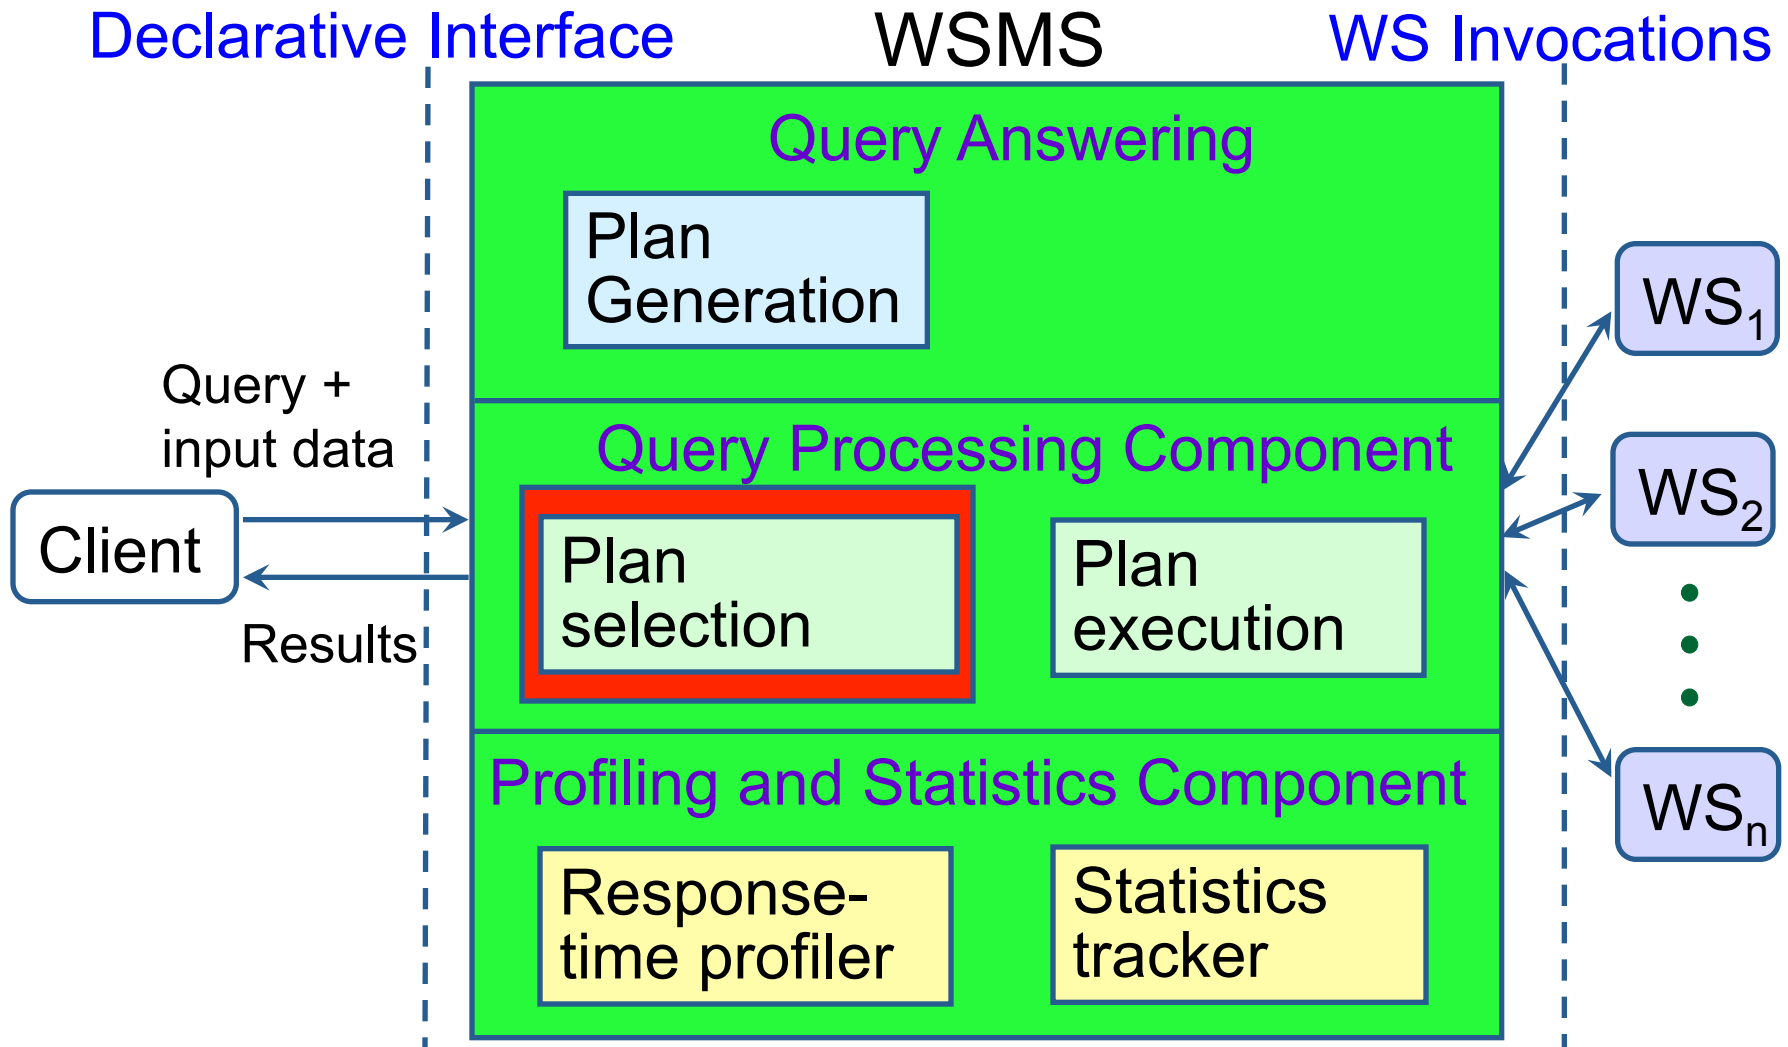


Web Services

Optimizing Web Service
Compositions

WSMS Architecture



Running Example

- Credit card company wants to send offers to people with:
 - a) credit rating > 600 , and
 - b) payment history = “good” on prior credit card
- Company has at its disposal:
 - L : List of potential recipients (identified by SSN)
 - WS_1 : SSN \rightarrow credit rating
 - WS_2 : SSN \rightarrow cc number(s)
 - WS_3 : cc number \rightarrow payment history


Web Services

- Company has at its disposal:
 - L : List of potential recipients (identified by SSN)
 $L^{\text{fff}}(\text{SSN}, \text{name}, \text{address})$
 - WS_1 : SSN \rightarrow credit rating
 $WS_1^{\text{bf}}(\text{SSN}, \text{CR})$
 - WS_2 : SSN \rightarrow cc number(s)
 $WS_2^{\text{bf}}(\text{SSN}, \text{CC})$
 - WS_3 : cc number \rightarrow payment history
 $WS_3^{\text{bf}}(\text{CC}, \text{PH})$


Query Evaluation

$L^{\text{ff}}(\text{SSN}, \text{name}, \text{address}), WS_1^{\text{bf}}(\text{SSN}, \text{CR}), WS_2^{\text{bf}}(\text{SSN}, \text{CC}),$
 $WS_3^{\text{bf}}(\text{CC}, \text{PH})$



Plan 1:



Plan2:



Which plan is better?



- Assume joins are “free” (cost 0)
- Cost metric: steady-state throughput

Plan 1 is never worse

Query Optimization Primer

- Possible query plans: P_1, \dots, P_n
- Data/access statistics: S
- Execution cost metric: $\text{cost}(P_i, S)$
- GOAL: Find least-cost plan

Query Optimization Primer

- Possible query plans: P_1, \dots, P_n
- Data/access statistics: S
- Execution cost metric: $\text{cost}(P_i, S)$
- GOAL: Find least-cost plan

Queries and Plans

Input: Evaluation plans using

- a set of extensional data sources L ,
- a set of web services WS_1, \dots, WS_n
- precedence constraints for WSs:

output of WS_i may be needed as input for WS_j

$$WS_i \rightarrow WS_j$$

E.g., $WS_2: SSN \rightarrow ccn$ and $WS_3: ccn \rightarrow ph$

$$WS_2 \rightarrow WS_3$$


→ Precedence DAG defines space of query plans

Query Optimization Primer

- Possible query plans: P_1, \dots, P_n
- Data/access statistics: S
- Execution cost metric: $\text{cost}(P_i, S)$
- GOAL: Find least-cost plan


Statistics

- 1) WS response times
- 2) WS selectivity



Statistics: Response Times

r_i : per-tuple response time of WS_i from client



- $r_i \approx 1/\text{throughput}$, can be reduced by batching, parallel calls
- Assume independent response times within query plans

Statistics: Selectivity

s_i : selectivity of WS_i

Average # output tuples per input tuple to WS_i
including post-filtering in query plan

WS_1 : SSN \rightarrow cr *then* filter cr > 600

If 90% of SSNs have cr > 600 then $s_1 = 0.9$

WS_2 : SSN \rightarrow ccn

If on average each SSN has 2 credit cards then $s_2 = 2.0$


Assumption: independent selectivities within query plans

Query Optimization Primer

- Possible query plans: P_1, \dots, P_n
- Data/access statistics: S
- Execution cost metric: $\text{cost}(P_i, S)$
- GOAL: Find least-cost plan

Bottleneck Cost Metric

Lunch Buffet



- Average *per-tuple* processing time:
response time of the slowest (bottleneck) stage in pipeline
- Observation: assume selectivity=1 in this example

Cost Equation for Plan P

$R_i(P)$: Predecessors of WS_i in plan P

- Fraction of input tuples seen by $WS_i = \prod_{j \in R_i(P)} s_j$
- WS_i response time per input tuple = $(\prod_{j \in R_i(P)} s_j) \cdot r_i$
- Bottleneck cost metric:


$$\text{cost}(P) = \max_{1 \leq i \leq n} \left(\left(\prod_{j \in R_i(P)} s_j \right) \cdot r_i \right)$$

(assumes WSMS processing is not the bottleneck)

Contrast with Sum Cost Metric

$$\text{cost}(P) = \sum_{1 \leq i \leq n} ((\prod_{j \in R_i(P)} s_j) \cdot r_i)$$

“Polite” Lunch Buffet




Problem Statement

- **Input:**
 - Web services WS_1, \dots, WS_n
 - Response times r_1, \dots, r_n
 - Selectivities s_1, \dots, s_n
 - Precedence constraints among web services
- **Output:**
 - Web services arranged into a plan P
 - P respects all precedence constraints
 - $\text{cost}(P)$ is minimized

No Precedence Constraints

Theorem: If all selectivities ≤ 1 , optimal to order linearly by r_i (selectivities are irrelevant)

General case (optimal):



With Precedence Constraints

$$\text{cost}(P) = \max_{1 \leq i \leq n} \left(\left(\prod_{j \in R_i(P)} s_j \right)^{\bullet r_i} \right)$$

With Precedence Constraints

$$\text{cost}(P) = \sum_{1 \leq i \leq n} ((\prod_{j \in R_i(P)} s_j) \cdot r_i)$$

Sum cost metric

- Hard to even obtain a factor $O(n^\theta)$ of optimal

With Precedence Constraints


$$\text{cost}(P) = \max_{1 \leq i \leq n} \left(\left(\prod_{j \in R_i(P)} s_j \right) \cdot r_i \right)$$

Bottleneck (max) cost metric


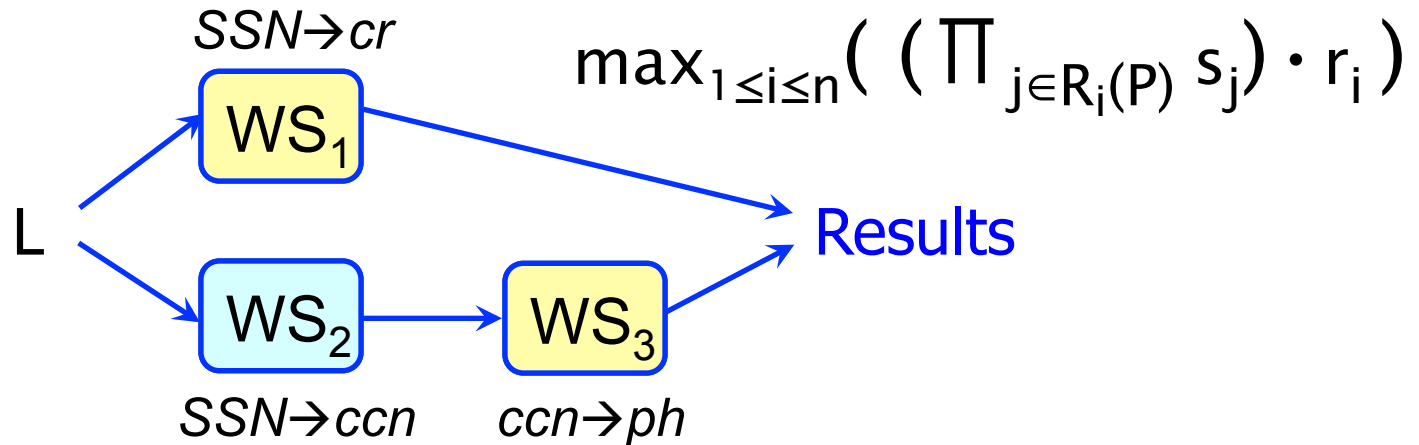
- Surprisingly, optimal solution in polynomial time
- $O(n^5)$ algorithm in the paper
 - Add one WS at a time to the plan
 - WS chosen by solving a linear program

Example Revisited


Plan 1:



Plan 2:



Selective



Proliferative



Precedence constraint